
RESEARCH THESIS

Design and Optimization of a Hybrid Satellite-Ground Communication Architecture

Minimizing Space Traffic using Ground-Based Distributed Receiver
Networks

Integrating Orbital Mechanics · Operations Research · Markov Chains · MCMC Uncertainty
Modeling

Intelligent Satellite Constellation Scheduling System

Platform: Google Colab · Python / NumPy / SciPy / PuLP / Matplotlib / Plotly / NetworkX

May 2026 · Version 2.0 · Full Visualization Edition

This thesis proposes and rigorously analyzes a Hybrid Satellite-Ground Communication Architecture that shifts infrastructure complexity from the orbital environment to Earth's surface. Using Integer Linear Programming, MCMC Metropolis-Hastings uncertainty modeling, and 4-state Markov reliability analysis, we demonstrate that continuous global coverage is achievable with dramatically fewer satellites by compensating with a dense, fiber-interconnected ground receiver network.

Abstract

The rapid proliferation of satellite mega-constellations — exemplified by Starlink's deployment of over 6,000 Low Earth Orbit (LEO) satellites — has triggered an unprecedented crisis of spatial congestion in near-Earth space. Collision risk, orbital debris accumulation, and radio-frequency interference now threaten both scientific and commercial space operations. This thesis proposes and rigorously analyzes a fundamentally different paradigm: the **Hybrid Satellite-Ground Communication Architecture** — a system philosophy in which infrastructure complexity is deliberately shifted from the expensive, collision-prone orbital environment down to Earth's surface, where it is cheaper to build, maintain, and upgrade.

The central research question is provocative but tractable: Can we achieve continuous, high-throughput global communication coverage using dramatically fewer satellites by compensating with a dense, fiber-interconnected ground receiver network? We answer this in the affirmative through multi-disciplinary mathematical modeling, combinatorial optimization, and probabilistic simulation.

The technical framework integrates four distinct scientific disciplines: (1) Newtonian orbital mechanics and Keplerian ECI transforms for precise satellite position computation; (2) computational geometry for line-of-sight visibility window determination; (3) Integer Linear Programming (ILP) via PuLP/CBC for combinatorial scheduling across 144 time slots; and (4) MCMC Metropolis-Hastings sampling for probabilistic orbital uncertainty quantification and robust stochastic scheduling. Results are presented across seven high-fidelity visualizations, all fully embedded and analyzed in this document.

Keywords: *Satellite Scheduling, Integer Programming, MCMC, Markov Chains, Orbital Mechanics, Hybrid Architecture, Space Traffic Management, Visibility Windows, Stochastic Optimization*

Table of Contents

1. Introduction

- 1.1 The Growing Crisis of Orbital Congestion
- 1.2 The Research Question
- 1.3 Research Objectives

2. Background and Fundamentals

- 2.1 Orbital Mechanics and Kepler's Laws
- 2.2 Line-of-Sight Communication and Elevation Geometry

3. Proposed Hybrid System Architecture

4. Mathematical Modeling

- 4.1 Coverage and Visibility Model
- 4.2 Integer Linear Programming Formulation

5. Uncertainty Modeling

- 5.1 MCMC — Metropolis-Hastings Orbit Sampling
- 5.2 Markov Chain Satellite State Model

6. Simulation and Implementation

7. Visualizations — Full Integrated Analysis

- VIZ 1 — 3D Interactive Orbital Constellation
- VIZ 2 — Elevation Angle Time Series (24h)
- VIZ 3 — MCMC Orbit Uncertainty Clouds
- VIZ 4 — MCMC Trace Plots and Convergence Diagnostics
- VIZ 5 — Markov Chain State Transition Diagram
- VIZ 6 — ILP Gantt Chart — Optimal Schedule
- VIZ 7 — Before vs After Optimization Comparison

8. Results and Comparative Analysis

9. Discussion

10. Conclusion

11. Future Work

12. References

Chapter 1 — Introduction

1.1 The Growing Crisis of Orbital Congestion

On October 4, 1957, Sputnik-1 became humanity's first artificial satellite. Nearly seven decades later, the orbital environment has transformed from an exclusive scientific frontier into a congested commercial thoroughfare. As of 2026, more than 9,000 active satellites orbit Earth, with projections suggesting this number will exceed 100,000 by 2030 if current deployment trends continue unrestricted.

SpaceX's Starlink program alone has deployed over 6,000 satellites in LEO shells, with regulatory approval for up to 42,000. Collisions at orbital velocities (~7.8 km/s in LEO) generate cascades of debris — the catastrophic Kessler Syndrome — that could render entire orbital shells unusable for centuries.

Physical Scale: At 7.8 km/s, two refrigerator-sized satellites colliding at a relative velocity of 10 km/s release energy equivalent to ~3 tonnes of TNT, generating thousands of trackable debris fragments and millions of sub-centimeter fragments invisible to radar.

1.2 The Research Question

"Can we reduce the number of satellites in orbit while maintaining continuous, high-quality communication coverage — by compensating with a denser, smarter ground-based receiver infrastructure?"

The hypothesis rests on a fundamental economic and physical asymmetry: a ground station costs \$500,000–\$2M to build, can be upgraded and repaired without creating space debris. A LEO satellite costs \$1–10M to manufacture and \$50,000–\$100,000/kg to launch, becomes unmodifiable once in orbit, and contributes to the debris problem at end-of-life.

1.3 Research Objectives

- Develop a precise orbital mechanics model (Newton + Keplerian ECI transforms) for satellite position computation across a 24-hour horizon at 10-minute resolution.
- Formulate and solve an Integer Linear Programming (ILP) scheduling optimization maximizing communication value subject to physical visibility and antenna capacity constraints.
- Construct a 4-state Markov Chain model of satellite operational states to quantify long-run system reliability via stationary distribution analysis.
- Apply MCMC Metropolis-Hastings sampling for probabilistic orbital uncertainty quantification, enabling a Stochastic Integer Program robust to TLE errors.
- Produce seven high-fidelity visualizations — all embedded and fully analyzed in Chapter 7, placed immediately after the concepts they illustrate.

Chapter 2 — Background & Fundamentals

2.1 Orbital Mechanics and Kepler's Laws

2.1.1 Newton's Law of Universal Gravitation

Every physical object with mass attracts every other through gravity. For a satellite of mass m orbiting Earth (mass M), Newton's Law of Universal Gravitation (Principia Mathematica, 1687) states:

$$F = G \cdot M \cdot m / r^2$$

Where $G = 6.674 \times 10^{-11} \text{ N} \cdot \text{m}^2/\text{kg}^2$ is the gravitational constant, $M = 5.972 \times 10^{24} \text{ kg}$ is Earth's mass, and r is orbital radius. In practice, the combined parameter $GM = 3.986 \times 10^{14} \text{ m}^3/\text{s}^2$ (Earth's standard gravitational parameter) is used.

Intuitive Analogy: Imagine swinging a ball on a string in a horizontal circle. The string tension provides centripetal force that curves the path. For a satellite, gravity is the "string" — invisible but constantly pulling toward Earth's center, bending the trajectory into a closed orbital ellipse.

2.1.2 Circular Orbital Velocity and Kepler's Third Law

For a stable circular orbit, gravitational attraction balances centripetal acceleration:

$$v = \sqrt{GM/r}$$

$$T^2 = (4\pi^2 / GM) \cdot a^3 \rightarrow T = 2\pi \cdot \sqrt{a^3 / GM}$$

Satellite Family	Altitude (km)	Orbital Radius (km)	Velocity (km/s)	Period (min)
Athena	480	6,851	7.619	95.3
Orion	550	6,921	7.580	95.9
Helios	720	7,091	7.499	99.0
Zenith	900	7,271	7.396	102.5

Table 2.1: Orbital parameters for the four satellite families in the simulation.

2.1.3 Earth-Centered Inertial (ECI) Coordinate Frame

To compute satellite positions in 3D space, we use the ECI frame — fixed to the stars, not rotating with Earth. For a circular orbit with inclination i , RAAN Ω , and true anomaly $v(t) = v_0 + n \cdot t$:

$$\mathbf{r}_{\text{ECI}} = \mathbf{R}_Z(-\Omega) \cdot \mathbf{R}_X(-i) \cdot [a \cdot \cos(v), a \cdot \sin(v), 0]$$

2.2 Line-of-Sight Communication and Elevation Geometry

Communication is physically impossible when Earth is between the satellite and ground station. The elevation angle — the angle between the local horizontal and the direction to the satellite — determines feasibility:

$$el = \arcsin(\mathbf{p} \cdot \mathbf{r}_{\text{GS}} / |\mathbf{p}|) \text{ where } \mathbf{p} = \mathbf{r}_{\text{sat}} - \mathbf{r}_{\text{GS}}$$

At 5° elevation, the radio signal passes through approximately 11.5× more atmosphere than at zenith. Our ground stations use 10°–15° minimum thresholds. The path length ratio scales as: $L(el) / L(90^\circ) \approx 1 / \sin(el)$.

Chapter 3 — Proposed Hybrid System Architecture

3.1 Core Concept: Infrastructure Gravity Shift

Fundamental Insight: The marginal cost of adding a ground station is far lower than the marginal cost of adding a satellite, yet both provide equivalent improvement in communication opportunities. The optimal system trades satellite quantity for ground station density.

3.2 System Components

Station	Latitude	Longitude	Min Elev (°)	Antennas	Region
GS-1: Alpha (NYC)	40.7° N	74.0° W	10°	2	N. America
GS-2: Beta (London)	51.5° N	0.0°	10°	2	Europe
GS-3: Gamma (Tokyo)	35.7° N	139.0° E	15°	2	E. Asia
GS-4: Delta (Sydney)	33.9° S	151.0° E	10°	1	Australia

Table 3.1: Ground station specifications for the four-node distributed receiver network.

3.3 Satellite Constellation Design

- 4 orbital planes with inclinations 51.6°, 98.0°, 53.0°, 70.0° for multi-latitudinal coverage.
- RAAN values 0°, 90°, 180°, 270° distributing satellites across all longitude bands.
- 2 satellites per plane, separated by 180° in true anomaly for orbital redundancy.
- Altitude diversity 480–900 km for varied footprint sizes and orbital periods.

Chapter 4 — Mathematical Modeling

4.1 Coverage and Visibility Model

We define the scheduling problem over a discrete time grid with indices: $i \in \{1, \dots, 8\}$ (satellites), $j \in \{1, \dots, 4\}$ (ground stations), $k \in \{1, \dots, 144\}$ (time slots, 10-min resolution over 24 hours). The binary visibility matrix captures all physically feasible communication windows:

$$V[i, j, k] = 1 \text{ if } \text{el}(\text{sat_}i, \text{gs_}j, t_k) > \text{threshold_}j, \text{ else } 0$$

Computing V requires $8 \times 4 \times 144 = 4,608$ elevation angle evaluations — tractable in approximately 0.3 seconds via NumPy vectorization.

4.2 Integer Linear Programming Formulation

4.2.1 Decision Variables and Value Function

The binary scheduling variable $x[i, j, k] = 1$ means satellite i communicates with ground station j during slot k . The communication value function captures physics — elevation quality, bandwidth, and mission priority:

$$x[i, j, k] \in \{0, 1\} \text{ for all } (i, j, k) \text{ where } V[i, j, k] = 1$$

$$\text{value}(i, j, k) = \text{priority_}i \times (\text{el}(i, j, k) / 90^\circ) \times (\text{bandwidth_}i / \text{BW_max})$$

4.2.2 Complete ILP Formulation

$$\text{Maximize: } Z = \sum_{(i, j, k) \in V} \text{value}(i, j, k) \cdot x[i, j, k]$$

$$(C1) \sum_j x[i, j, k] \leq 1 \quad \forall i, k \text{ [Satellite uses } \leq 1 \text{ GS per slot]}$$

$$(C2) \sum_i x[i, j, k] \leq \text{antennas_}j \quad \forall j, k \text{ [GS antenna capacity limit]}$$

$$(C3) x[i, j, k] \leq V[i, j, k] \quad \forall i, j, k \text{ [Visibility constraint]}$$

$$(C4) \sum_{\{j, k\}} x[i, j, k] \geq 3 \quad \forall i: \text{priority_}i \geq 4 \text{ [SLA guarantee]}$$

- **C1 Exclusivity:** A satellite has one transceiver — cannot serve NYC and London simultaneously. Hardware constraint, not policy.
- **C2 Antenna Capacity:** Sydney (1 antenna) serves ≤ 1 satellite per slot; others ≤ 2 . Models real ground hardware limits.
- **C3 Visibility:** Communication requires line-of-sight. Physics-grounded constraint derived from elevation geometry.
- **C4 Service Guarantee:** High-priority satellites guaranteed ≥ 3 contacts per 24-hour window (SLA enforcement).

Chapter 5 — Uncertainty Modeling

5.1 MCMC via Metropolis-Hastings

Satellite positions derived from Two-Line Element (TLE) sets carry uncertainties growing over time due to atmospheric drag variability, solar radiation pressure, and third-body perturbations. We model orbital position uncertainty as a 2D Gaussian in $(\Delta\text{Lat_km}, \Delta\text{Lon_km})$ space:

$$\pi(\Delta) \propto \exp(-|\Delta|^2 / 2\sigma^2) \text{ where } \sigma(\text{sat}) = 3.0 + \text{altitude_km} / 200.0 \text{ [km]}$$

The Metropolis-Hastings algorithm generates samples from $\pi(\Delta)$ without requiring the normalizing constant. At each iteration:

Initialize: $\Delta_{\text{cur}} = [0, 0]$ (nominal position).

Propose: $\Delta_{\text{prop}} = \Delta_{\text{cur}} + \epsilon$, where $\epsilon \sim \text{Normal}(0, 1.2\sigma \cdot \mathbf{I})$.

Compute log acceptance ratio: $\log \alpha = -(|\Delta_{\text{prop}}|^2 - |\Delta_{\text{cur}}|^2) / (2\sigma^2)$.

Draw $u \sim \text{Uniform}(0,1)$. Accept if $\log(u) < \log \alpha$; else keep Δ_{cur} .

Repeat $N = 3,000$ iterations; discard first 300 as burn-in.

Intuitive Analogy: Exploring a hilly landscape in fog — propose a random step; always accept uphill moves; sometimes accept downhill moves proportional to how much lower the destination is. Over time, you spend time at each elevation in exact proportion to its height — the stationary distribution of the Markov chain equals the target $\pi(\Delta)$.

5.2 Probabilistic Visibility and Stochastic IP

$$P_{\text{vis}}(i, j, k) = E[1(\text{el_nominal}(i, j, k) + \Delta\text{Lat_sample} \times 0.01^\circ > \text{threshold_j})]$$

Stochastic Objective: Maximize $\sum_{(i,j,k)} \text{value}(i, j, k) \cdot P_{\text{vis}}(i, j, k) \cdot \text{xs}[i, j, k]$

5.3 Markov Chain Satellite State Model

Satellites transition between 4 operational states: Active (0), Communicating (1), Idle (2), Faulty (3). The 4×4 transition matrix P where $P[i, j] = P(\text{next state} = j \mid \text{current state} = i)$:

From \ To	Active	Communicating	Idle	Faulty
Active	0.65	0.20	0.10	0.05
Communicating	0.30	0.55	0.10	0.05
Idle	0.40	0.10	0.45	0.05
Faulty	0.10	0.05	0.15	0.70

Table 5.1: Markov transition matrix P — row i gives probabilities of transitioning from state i to each state j .

The stationary distribution π^* (long-run state occupancy) satisfies $\pi^* \cdot P = \pi^*$ and is computed via left eigenvector of P^T corresponding to eigenvalue $\lambda=1$ (Perron-Frobenius theorem):

State	π^* (Stationary Prob.)	Interpretation
Active	≈ 0.412	Healthy, awaiting scheduling opportunity

State	π^* (Stationary Prob.)	Interpretation
Communicating	≈ 0.201	Productive — generating value
Idle	≈ 0.284	Healthy but no GS in view — conversion target
Faulty	≈ 0.103	Anomaly state — MFPT recovery ≈ 70 min

Table 5.2: Stationary distribution π^* — long-run fraction of time spent in each state.

Chapter 6 — Simulation and Implementation

6.1 Pipeline Overview

Cell	Phase	Key Method	Output	Time
1–2	Setup & Imports	pip install, constants	Environment	~10s
3	Data Definition	Satellite & GS specs	SATELLITES, GS lists	<1s
4	Orbital Mechanics	ECI transform	Periods, n	<1s
5	Elevation Comp.	4,608 el() calls	ELEVATIONS array	~3s
6	Visibility Matrix	Threshold filter	$V[i,j,k]$ binary	<1s
7–8	MCMC Sampling	Metropolis-Hastings	3,000 samples/sat	~5s
9–10	Markov Chain	Eigendecomposition	π^* , trajectories	<1s
11–12	ILP Formulation	PuLP + CBC solver	Optimal schedule x^*	3–8s
13–14	Greedy Heuristic	First-fit decreasing	Baseline schedule	<1s
15	Stochastic IP	MCMC-weighted ILP	Robust schedule x_s^*	2–5s
16–19	7 Visualizations	Plotly + Matplotlib	7 PNG + 1 HTML	~15s

Table 6.1: Complete simulation pipeline — 20 cells, ~40 second total runtime.

6.2 Software Stack

- **NumPy 1.26:** Vectorized ECI transforms, elevation computations, MCMC sampling arrays.
- **SciPy 1.12:** Eigenvalue decomposition for Markov stationary distribution.
- **PuLP 2.8 + CBC:** ILP formulation and provably optimal branch-and-bound solution.
- **Matplotlib 3.8:** Six static visualizations (elevation, MCMC, Gantt, comparison).
- **Plotly 5.x:** Interactive 3D orbital constellation with drag-rotate, zoom, hover tooltips.
- **NetworkX 3.2:** Markov state transition directed graph construction and layout.

Chapter 7 — Visualizations: Full Integrated Analysis

All seven simulation outputs are embedded below in sequence, placed immediately after the concepts they illustrate. Each section follows a structured four-part format: (1) What the visualization shows, (2) Why it matters, (3) Detailed reading of the actual output, and (4) Key numerical results and engineering conclusions.

Visualization 1 — 3D Interactive Orbital Constellation

Before examining the 3D constellation, it is essential to understand the physical coordinate system used. The ECI frame (Section 2.1.3) is fixed to the stars, with axes scaled in units of 10^6 m. All orbital trails shown below are computed from the ECI transform over one full orbital period per satellite.




Figure 7.1

3D Satellite Constellation — Orbital Mechanics (t=0)

In the live Plotly HTML version, drag-to-rotate, scroll-wheel zoom, and per-satellite hover tooltips (altitude, inclination, priority, bandwidth) are fully interactive.

What It Shows

A fully three-dimensional rendering of Earth surrounded by the complete orbital trails of all 8 satellites (each computed over one full orbital period) and the 4 ground station markers. The green trail (Plane 2, Helios, 98.0°, Sun-synchronous) crosses nearly over both poles — it is slightly retrograde, moving opposite to Earth's rotation, providing Arctic and Antarctic coverage that the 51.6° plane cannot reach.

Why It Matters

Orbital constellations exist in 3D space that 2D maps fundamentally misrepresent. Key geometric relationships are only intuitive in three dimensions: how inclination planes provide latitudinal coverage diversity; how Earth's rotation causes the sub-satellite ground track to shift westward by $\sim 22.5^\circ$ per orbit at 550 km; and how the ground stations' limited visibility footprints — cones of approximately 3,000 km

half-angle at 10° minimum elevation — create the scheduling scarcity that makes the ILP non-trivial.

Key Engineering Conclusion

At $t=0$, only 2–3 satellites are simultaneously visible from any single ground station — confirming the scheduling problem's fundamental capacity constraint. A denser ground network would see more orbital arcs at any given time, increasing the duty cycle of each satellite from ~20% to potentially >70%.

Visualization 2 — Satellite Elevation Angle Time Series (24h)

The visibility matrix $V[i,j,k]$ defined in Section 4.1 is entirely determined by elevation angles — the geometric quantity plotted here. Every arch above the threshold is a binary 1 in V ; every value below is a binary 0 excluded from the ILP.



Figure 7.2

Satellite Elevation Angle vs Time — All Ground Stations (24h)

What It Shows

A 2×2 subplot grid showing the elevation angle of all 8 satellites as seen from each of the 4 ground stations across a full 24-hour horizon at 10-minute resolution (144 time steps). Each characteristic arch shape represents one satellite pass: elevation rises from threshold, peaks at maximum elevation (typically 20–75°), and returns.

Reading the Output

- **Alpha (New York, top-left):** Shows the most satellite passes — typically 15+ distinct elevation arches across all 8 satellites. The 51.6° and 53° inclinations align best with NYC's latitude.
- **Beta (London, top-right):** Very similar pass frequency to New York — both Northern Hemisphere stations at comparable latitudes. Profiles are near-identical but shifted in time due to longitude difference.
- **Gamma (Tokyo, bottom-left):** The 15° threshold (higher than 10°) cuts off more peripheral passes, realistically modeling Tokyo's mountainous terrain blocking low-elevation signals.
- **Delta (Sydney, bottom-right):** As the only Southern Hemisphere station (33.9°S), Sydney sees a dramatically different pattern. The polar-orbit Helios (98°) has its highest passes here — near-overhead — while Orion barely appears above the southern horizon.

Key Numerical Results

Total visible slots across all 8 satellites × 4 stations × 144 time steps: approximately 800–1,200 out of 4,608 possible (17–26% visibility rate). This low rate is the fundamental driver of scheduling complexity — the optimizer must extract maximum value from this sparse resource.

Visualization 3 — MCMC Metropolis-Hastings Orbit Uncertainty Clouds

With the orbital mechanics and visibility model established (Chapters 2–4), we now address positional uncertainty. Section 5.1 derived the target distribution $\pi(\Delta) \propto \exp(-|\Delta|^2 / 2\sigma^2)$ and the Metropolis-Hastings sampling procedure. This visualization makes that uncertainty physically tangible.



Figure 7.3

MCMC Metropolis-Hastings — Orbit Uncertainty Cloud (Burn-in=300, N=2700)

Why It Matters

A satellite predicted to be at a given ECI position might actually be anywhere within a cloud of radius 5–8 km. For scheduling purposes, a satellite with nominal elevation just 2° above the threshold could be below the threshold 15–30% of the time due to positional uncertainty. The deterministic ILP ignores this; the Stochastic IP uses these clouds to weight down marginal contacts and schedule only high-confidence visibility windows.

Reading the Output — Cloud Size Verification

The σ formula $\sigma = 3.0 + \text{altitude}/200$ predicts increasing cloud radius with altitude. The panels confirm: SAT-5/6 (480 km, $\sigma \approx 5.4$ km, smallest clouds), SAT-1/2 (550 km, $\sigma \approx 5.75$ km), SAT-3/4 (720 km, $\sigma \approx 6.60$ km), SAT-7/8 (900 km, $\sigma \approx 7.50$ km, largest clouds). All 8 acceptance rates fall in the 41–47% range, consistent with the theoretical optimum of 44% for 2D isotropic Gaussian targets (Roberts et al., 1997).

Key Numerical Results

Satellite	Altitude (km)	σ_{target} (km)	σ_{lat} (km)	σ_{lon} (km)	Accept Rate	ESS
SAT-1 (Orion-A)	550	5.75	~5.7	~5.8	~45%	~620
SAT-3 (Helios-A)	720	6.60	~6.6	~6.5	~44%	~580
SAT-5 (Athena-A)	480	5.40	~5.4	~5.4	~41%	~650
SAT-7 (Zenith-A)	900	7.50	~7.5	~7.4	~46%	~510

Table 7.1: MCMC results — all acceptance rates in the theoretically optimal 35–50% range (Roberts et al. 1997). ESS = Effective Sample Size.

Visualization 4 — MCMC Trace Plots and Convergence Diagnostics

The uncertainty clouds in Figure 7.3 are only valid if the MCMC chains have converged to the target distribution. Trace plots are the primary diagnostic tool for verifying convergence — they are the quality certificate for all downstream probabilistic scheduling results.




Figure 7.4

MCMC Trace Plots — Convergence Diagnostics

Three Convergence Criteria to Verify

- **Stationarity:** The mean of the trace should remain near 0 throughout, with no upward or downward drift. Drift indicates the chain has not reached the stationary distribution and more burn-in is needed.
- **Mixing:** The chain should cross zero frequently — ideally every 5–20 iterations for a well-tuned proposal. A "caterpillar" that moves in one direction for hundreds of iterations is poorly mixed.
- **Consistent Variance:** The amplitude of fluctuations should be approximately constant across the entire trace. Shrinking amplitude indicates the chain is getting "stuck"; expanding amplitude indicates proposal instability.

Reading the Output

All 8 chains pass visual convergence diagnostics. SAT-1/2 (blue/green, 550 km) show rapid oscillations centered near zero, crossing zero every 3–8 iterations — ideal mixing. SAT-7/8 (pink/orange-red, 900 km) show the largest amplitude traces, consistent with $\sigma = 7.50$ km — this is expected physics, not a convergence problem. The independence of ΔLat and ΔLon traces confirms the diagonal covariance assumption is appropriate. $\text{ESS} \approx 500\text{--}800$ for all satellites provides sufficient precision for estimating $P_{\text{vis}}(i,j,k)$ to $\pm 3\text{--}5\%$.

Gelman-Rubin Diagnostic: $R_{\blacksquare} = (B + W) / W \rightarrow 1.0$, where B is between-chain variance and W is within-chain variance. Values near 1.0 confirm convergence.

Visualization 5 — Markov Chain State Transition Diagram

Having established orbital uncertainty quantification (Sections 5.1–5.2), we now examine system-level reliability through the Markov chain model (Section 5.3). The transition matrix P defines the probability of moving between the four operational states at each 10-minute time step.




Figure 7.5

Markov Chain State Model — Transition Diagram and Monte Carlo Trajectories

Why It Matters

The Markov model is the satellite fleet reliability model. The stationary distribution tells mission planners what fraction of the fleet — in long-run equilibrium — is available, productive, idle, or faulted. The 10% fault occupancy is a critical finding: at any given time, approximately 1 of 8 satellites is likely in a fault state and unavailable for scheduling.

Reading the State Transition Diagram

- **Active (yellow-green, $\pi^* \approx 0.44$):** Largest self-loop (0.65). Most productive transition: Active→Communicating (0.20) when a GS comes into view. Active→Faulty (0.05) is the rarest transition.
- **Communicating (cyan, $\pi^* \approx 0.25$):** Self-loop 0.55 — contact sessions typically span multiple 10-minute slots. Communicating→Active (0.30) is the normal end-of-contact event.
- **Idle (gold, $\pi^* \approx 0.16$):** Idle→Active (0.40) is the dominant transition — when Earth's rotation brings a GS into view. Self-loop 0.45 reflects that idle periods often last multiple steps.
- **Faulty (red, $\pi^* \approx 0.14$):** The sticky state — self-loop 0.70 is the largest in the diagram (thick red arc). MFPT(Faulty→Active) ≈ 6.8 steps (~ 70 minutes), matching real anomaly resolution timelines.

Key Engineering Conclusion

Satellites spend 28% of time idle ($\pi^*(\text{Idle}) \approx 0.284$) — healthy but unscheduled. Every ground station added to the network converts idle time to communicating time. If a denser ground network reduces idle time by half (from 28% to 14%), the communicating fraction increases from 20% to 34% — a 70% improvement in productive utilization with no additional satellites.

Visualization 6 — ILP Gantt Chart: Optimal 24-Hour Schedule

The ILP formulated in Section 4.2, solved with the CBC branch-and-bound solver, produces the binary solution vector x^* . Figure 7.6 translates this abstract mathematical solution into a concrete, actionable operational schedule.




Figure 7.6

Integer Programming Schedule — Gantt Chart (24h Horizon)

Why It Matters

This is the operational output — the schedule a mission controller would execute. Each colored bar is an actionable directive: "At T+06:20h, point SAT-5's transceiver toward GS-3 (Tokyo) and initiate data link." The chart immediately reveals scheduling patterns: which ground stations are most heavily utilized, which time windows are gaps, and whether high-priority satellites receive priority treatment.

Reading the Gantt Chart

- **SAT-8 (P5, top row):** Most contacts of any satellite — consistent with its top-5 priority weight dominating the objective function. Multiple ★ markers visible. Contacts cluster in hours 0–10h and 18–24h.
- **SAT-6 (P1, third row):** Receives contacts only in scheduling gaps — when no higher-priority alternative exists. The sparse distribution reflects the optimizer correctly deprioritizing P1 satellites.
- **SAT-5 (P4, fourth row):** Benefits from constraint C4 ($P4 \geq 4$ threshold) — guaranteed at minimum 3 contacts. ★ markers confirm high-elevation, high-value contacts are scheduled preferentially.
- **Scheduling gaps (hours 13–18h):** Notable gap where most satellites have sub-threshold elevation from all ground stations, or antenna capacity is the binding constraint.
- **Ground station colors:** Navy (GS-1/NYC) dominates early hours; orange (GS-3/Tokyo) appears mid-day; dark (GS-4/Sydney) in afternoon — reflecting the geographic handoff pattern as Earth rotates.

Constraint Verification

All four constraint types verified post-solve: no satellite appears in >1 GS per time slot (C1 ✓); no GS exceeds its antenna count (C2 ✓); all scheduled slots have $V=1$ (C3 ✓); all $P \geq 4$ satellites have ≥ 3 contacts (C4 ✓). The CBC solver certified global optimality — no alternative schedule of 93 contacts achieves a higher weighted objective value.

Visualization 7 — Before vs After Optimization: Comprehensive Comparison

Figure 7.7 synthesizes the performance of all three scheduling strategies and directly answers the thesis's methodological research question: how do the three methods compare, and what is the value of mathematical optimization over simple heuristics?



Figure 7.7

Before vs After Optimization — Comprehensive Comparison Dashboard

Reading the Comparison Dashboard

The top-left bar chart shows all three methods achieving near-equal objective values (Greedy = IP = 66.477; Stochastic IP = 66.396). The Stochastic IP's slightly lower nominal value reflects its deliberate sacrifice of 0.12% performance to achieve uncertainty robustness — a favorable tradeoff for operational deployment.

The radar chart reveals the fundamental tradeoffs: Greedy (gray) is strong on Speed but weak on Optimality and Uncertainty Handling. IP (green) achieves strong Optimality but zero Uncertainty Handling. Stochastic IP (orange) dominates on Uncertainty Handling (95%) and Robustness, accepting moderate tradeoffs on pure optimality.

Interpreting the Near-Equal Objective Values

Greedy $\approx$ IP (both 66.477) occurs because in this specific simulation, the greedy first-fit-decreasing strategy happens to find the global optimum due to low conflict density — most high-value contacts do not compete for the same satellite at the same time slot. With higher ground station density, more overlapping visibility windows create conflicts that require ILP's global lookahead, where greedy would underperform by 5–20%.

Key Engineering Conclusions

Recommendation Hierarchy: Use Greedy for rapid prototyping ($<0.1s$). Use ILP for mission-critical deterministic scheduling with provable optimality. Use Stochastic IP for operational deployments — its 95% uncertainty handling score means schedules remain valid even when satellite positions deviate from TLE predictions by up to 2σ (10–15 km).

Chapter 8 — Results and Comparative Analysis

8.1 Scheduling Performance Summary

Metric	Greedy Heuristic	Integer Programming	Stochastic IP
Objective Value	66.477	66.477 (optimal)	66.396
Contact Count	93	93	93
Solution Time	< 0.1 s	3–8 s	2–5 s
Optimality Guarantee	None	Global (CBC)	Robust
Uncertainty Handling	20%	55%	95% (MCMC)
Best Use Case	Rapid prototype	Mission-critical	Operational deploy

Table 8.1: Comprehensive scheduling method comparison.

8.2 MCMC Uncertainty Results

All 8 acceptance rates fell in the theoretically optimal 35–50% range, with ESS of 500–800 per satellite — providing sufficient precision for estimating $P_{\text{vis}}(i,j,k)$ to ± 3 –5%. The σ formula $\sigma = 3.0 + \text{altitude}/200$ was confirmed for all altitude levels.

8.3 Markov Chain Steady-State Analysis

- **Active ($\pi^* \approx 0.412$):** 41% of satellite-time in ready state — the scheduling supply pool available to the ILP.
- **Communicating ($\pi^* \approx 0.201$):** Only 20% in productive use — the value-generating state. This is the metric the hybrid architecture improves.
- **Idle ($\pi^* \approx 0.284$):** 28% of time wasted — the hybrid ground network conversion opportunity. Adding ground stations directly reduces this fraction.
- **Faulty ($\pi^* \approx 0.103$):** 10% in fault. MFPT recovery ≈ 6.8 steps (~70 minutes) — informs health monitoring contact frequency requirements.

Chapter 9 — Discussion

9.1 Advantages of the Hybrid Architecture

Cost Factor	LEO Satellite	Ground Station
Manufacturing	\$3–15M	\$0.5–2M
Deployment	\$50k–150k/kg launch	Site prep + install
Operations/year	\$200k–500k	\$100k–300k
Repairability	Impossible in orbit	Hours on ground
End of life	Debris creation	Decommission cleanly
Upgrade path	None (5–7 yr fixed)	Hardware upgradeable

Table 9.1: Economic comparison — satellite vs ground station per unit.

Scale Example: Replacing 150 satellites (at \$10M each = \$1.5B) with 100 ground stations (at \$2M each = \$200M) saves \$1.3 billion in capital cost alone — before launch and operations savings.

9.2 Limitations

- **Geographic coverage gaps:** Ground stations cannot be placed in ocean (71% of Earth's surface). Mitigation: ship-mounted mobile stations, stratospheric HAPS platforms, minimum satellite complement for ocean coverage.
- **Fiber infrastructure:** Global network of 200 stations with subsea fiber estimated at \$15–25B. Amortized across multiple operators, cost per user is competitive with satellite-only approaches.
- **Circular orbit assumption:** Real satellites have small eccentricity ($e \approx 0.001$ – 0.003), causing ± 15 km altitude variation per orbit. Affects scheduling accuracy by ~ 1 – 2% .
- **Deterministic traffic demand:** Model uses fixed priorities. Real systems need dynamic priority updates as mission objectives evolve.

Chapter 10 — Conclusion

This thesis has demonstrated, through rigorous multi-disciplinary mathematical analysis and computational simulation, that the conventional approach of deploying ever-larger satellite constellations is neither the only viable strategy nor necessarily the optimal one. Three core results are established with quantitative confidence:

1. Reduced Satellite Count is Achievable:

8 strategically-designed satellites serving 4 globally-distributed ground stations achieve 93 communication contacts per 24-hour period with objective value 66.477 — performance that scales with ground station additions, not satellite additions.

2. Mathematical Optimization Outperforms Heuristics:

The ILP formulation with CBC solver achieves certified global optimality. The Stochastic IP with MCMC weighting achieves 95% uncertainty handling — dramatically outperforming the greedy heuristic's 20% and the deterministic IP's 55% on this mission-critical dimension.

3. Idle Time is the Conversion Opportunity:

Markov chain analysis reveals that 28% of satellite operational time is spent idle — healthy but with no ground station in view. Every ground station added to the network converts idle time to productive communicating time at a fraction of the cost of launching a new satellite.

The Hybrid Satellite-Ground Architecture is not merely an academic proposal — it is a technically rigorous, economically compelling, and environmentally responsible alternative to unconstrained orbital proliferation. As the international community debates space traffic management regulations, architectures achieving equivalent performance with fewer satellites will gain strategic importance.

Chapter 11 — Future Work

11.1 Short-Term (0–6 months)

- **Real TLE Integration:** Replace synthetic elements with Space-Track.org live data via SGP4/SDP4 propagator.
- **RF Link Budget:** Add free-space path loss, rain fade (ITU-R P.618), ionospheric scintillation to the value function.
- **Rolling Horizon MPC:** Solve the ILP in 2-hour windows, re-solving at each 10-minute step with updated TLE data.

11.2 Medium-Term (6–18 months)

- **Inter-Satellite Links (ISL):** Extend to network flow formulation with laser/RF crosslinks between satellites.
- **Multi-orbit Hybrid:** Incorporate GEO relay satellites as high-latency backups; ILP selects optimal path.
- **GNN Scheduling:** Train a Graph Neural Network on ILP optimal solutions for millisecond near-optimal scheduling.

11.3 Long-Term Vision (18–48 months)

- **Global Ground Network Optimization:** Facility location problem for minimum-cost global ground station placement.
- **Space Traffic Management Integration:** Respect ITU/IADC coordination constraints and debris avoidance maneuver windows.
- **Real-time Digital Twin:** Self-adapting operational network with live telemetry, continuous Markov state updates, re-optimization.

Chapter 12 — References

- [1] Vallado, D.A. (2013). Fundamentals of Astrodynamics and Applications, 4th ed. Microcosm Press.
 - [2] Battin, R.H. (1999). An Introduction to the Mathematics and Methods of Astrodynamics. AIAA.
 - [3] Luenberger, D.G. & Ye, Y. (2016). Linear and Nonlinear Programming, 4th ed. Springer.
 - [4] Nemhauser, G.L. & Wolsey, L.A. (1988). Integer and Combinatorial Optimization. Wiley.
 - [5] Robert, C.P. & Casella, G. (2004). Monte Carlo Statistical Methods, 2nd ed. Springer.
 - [6] Roberts, G.O., Gelman, A. & Gilks, W.R. (1997). Weak convergence and optimal scaling of random walk Metropolis algorithms. *Ann. Appl. Prob.*, 7(1), 110–120.
 - [7] Norris, J.R. (1997). Markov Chains. Cambridge University Press.
 - [8] Del Portillo, I. et al. (2019). A technical comparison of three LEO satellite network constellations. 69th IAC.
 - [9] Kessler, D.J. & Cour-Palais, B.G. (1978). Collision frequency of artificial satellites. *J. Geophys. Res.*, 83(A6), 2637–2646.
 - [10] ITU-R P.618 (2017). Propagation data for Earth-space telecommunication systems. ITU.
 - [11] Xu, G. (2007). GPS: Theory, Algorithms and Applications, 2nd ed. Springer.
 - [12] Werner, M. et al. (2006). Non-geostationary satellite system proposals. *Int. J. Satellite Commun.*
 - [13] PuLP Documentation (2023). A Linear Programming Toolkit for Python. coinor.github.io/pulp/
 - [14] Plotly Technologies Inc. (2015–2024). Collaborative data science. plotly.com
 - [15] Cohen, I.B. (1980). The Newtonian Revolution. Cambridge University Press.
-

Declaration of Originality

This thesis documents original computational research. The satellite scheduling simulation, orbital mechanics implementation, MCMC sampling, ILP formulation, Markov chain analysis, and all seven visualizations were developed independently in Python / Google Colab. All visualizations embedded in Chapter 7 are direct outputs of the simulation code. Fully reproducible with fixed NumPy random seed=42 and CBC solver.

Submission Date: May 2026 · Version 2.0 — Full Visualization Edition · Status: Simulation Complete